

**VIVEKANAND EDUCATION SOCIETY'S
INSTITUTE OF TECHNOLOGY**

Department of Computer Engineering



Report on
**Managing Container Logs and build an ELK service
using Docker**

**Faculty Incharge
Richard Joesph**

Submitted by
Group No. 5
Ishaan Malhi (D17A, Roll no : 39)
Manish Manghwani (D17A, Roll No : 41)
Deepika Khatri (D17A, Roll no : 31)

(2016-17)

Table of Contents

Sr. No	Title	Page
1	Abstract	3
2	Introduction	4
3	Implementation	8
4	Results and Graphs	20
5	Conclusion	23
6	References	24

1. Abstract

Logging is one of the essential parts in any setup because logs are extremely useful when it comes to troubleshooting problems happening with your production app. When you decide to move your apps to the docker-managed environments, we realise that some of the approaches we were using don't work anymore — now the logs are staying inside the container and we need to find a way to make them persistent. It goes without saying that the easiest way is to put the logs to the volume but what if we have many instances of the app? In this case we have to use a centralised storage for logs.

We've decided to store our logs in one place using the ELK stack. ELK is a centralised log server and log management web interface which includes Elasticsearch (as a storage and a search engine), Logstash (a tool for log processing), and Kibana (UI for the Elasticsearch).

2. Introduction

What is Docker?

Docker is the world's leading software container platform. Developers use Docker to eliminate “works on my machine” problems when collaborating on code with co-workers. Operators use Docker to run and manage apps side-by-side in isolated containers to get better compute density. Enterprises use Docker to build agile software delivery pipelines to ship new features faster, more securely and with confidence for both Linux and Windows Server apps.

What is a Container?

A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings. Available for both Linux and Windows based apps, containerized software will always run the same, regardless of the environment. Containers isolate software from its surroundings, for example differences between development and staging environments and help reduce conflicts between teams running different software on the same infrastructure. Using containers, everything required to make a piece of software run is packaged into isolated containers. Unlike VMs, containers do not bundle a full operating system - only libraries and settings required to make the software work are needed. This makes for efficient, lightweight, self-contained systems and guarantees that software will always run the same, regardless of where it's deployed.

Containers and virtual machines have similar resource isolation and allocation benefits, but function differently because containers virtualize the operating system instead of hardware, containers are more portable and efficient.

Docker For Developers

Docker automates the repetitive tasks of setting up and configuring development environments so that developers can focus on what matters: building great software.

Developers using Docker don't have to install and configure complex databases nor worry about switching between incompatible language toolchain versions. When an app is dockerized, that complexity is pushed into containers that are easily built, shared and run. Onboarding a co-worker to a new codebase no longer means hours spent installing software and explaining setup procedures. Code that ships with Dockerfiles is simpler to work on: Dependencies are pulled as neatly packaged Docker images and anyone with Docker and an editor installed can build and debug the app in minutes.

What is the ELK Stack?

The ELK stack consists of Elasticsearch, Logstash, and Kibana. Although they've all been built to work exceptionally well together, each one is a separate project that is driven by the open-source vendor Elastic—which itself began as an enterprise search platform vendor. It has now become a full-service analytics software company, mainly because of the success of the ELK stack. Wide adoption of Elasticsearch for analytics has been the main driver of its popularity.

Data constantly flows into your systems, but it can quickly grow to be large and stale. As your data set grows larger, your analytics will slow up, resulting in sluggish insights. And this is likely to be a serious business problem. So, the BIG question for your big data is: how can you maintain valuable business insights?

Not long ago, an epiphany ran through the industry: analytics is, in essence, a search problem that needs coupling with good visualizations. So, there was a marriage: Lucene with all its search goodness was brought together with the distributed-computing goodness that is Elasticsearch. Logstash came onto the scene to normalize all kinds of time-series data. Pop in Kibana's ultra-simple visualization tool, and you have a complete analytics tool that can rival very expensive and scalable solutions from Oracle, Palantir, Tableau, Splunk, Microsoft, and others.

The ELK stack makes it much easier -- and much faster -- to search and analyze large data sets as your data grows larger and larger.

We should mention that ELK is quite versatile. Use the stack as a stand-alone application, or integrate with your existing applications to get the most current data. With Elasticsearch, you get all the features to make real-time decisions-all the time. You can use each of these tools separately, or with other products. For example, Kibana often goes together with Solr/Lucene. Although none of these is a project of the Apache Foundation, each part of the stack falls under the Apache 2 License. Elasticsearch owns both the intellectual property and the trademarks.

- **Elasticsearch — The Log Search Tool**

Elasticsearch is a juggernaut solution for your data extraction problems. A single developer can use it to find the high-value needles underneath all of your data haystacks.

- **Real-time data and real-time analytics.** The ELK stack gives you the power of real-time data insights, with the ability to perform super-fast data extractions from virtually all structured or unstructured data sources. Real-time extraction, and real-time analytics. Elasticsearch is the engine that gives you both the power and the speed.
- **Scalable, high-availability, multi-tenant.** With Elasticsearch, you can start small and expand it along with your business growth-when you are ready. It is built to scale horizontally out of the box. As you need more capacity, simply add another node and let the cluster reorganize itself to accommodate and exploit the extra hardware. Elasticsearch clusters are resilient, since they automatically detect and remove node failures. You can set up multiple indices and query each of them independently or in combination.
- **Full text search.** Under the cover, Elasticsearch uses Lucene to provide the most powerful full-text search capabilities available in any open-source product. The search features come with multi-language support, an extensive query language, geolocation support, and context-sensitive suggestions, and autocomplete.
- **Document orientation.** You can store complex, real-world entities in Elasticsearch as structured JSON documents. All fields have a default index, and you can use all the indices in a single query to get precise results in the blink of an eye.

Logstash — Routing Your Log Data

Logstash is a tool for log data intake, processing, and output. This includes virtually any type of log that you manage: system logs, webserver logs, error logs, and app logs. As administrators,

we know how much time can be spent normalizing data from disparate data sources. We know, for example, how widely Apache logs differ from NGINX logs.

Rather than normalizing with time-taking ETL (Extract, Transform, and Load), we recommend that you switch over to the fast track. Instead, you could spend much less time training Logstash to normalize the data, getting Elasticsearch to process the data, and then visualize it with Kibana. With Logstash, it's super easy to take all those logs and store them in a central location. The only prerequisite is a Java runtime, and it takes just two commands to get Logstash up and running. Using Elasticsearch as a backend datastore and Kibana as a frontend dashboard, Logstash will serve as the workhorse for storage, querying and analysis of your logs. Since it has an arsenal of ready-made inputs, filters, codecs, and outputs, you can grab hold of a very powerful feature-set with a very little effort on your part.

Think of Logstash as a pipeline for event processing: it takes precious little time to choose the inputs, configure the filters, and extract the *relevant, high-value* data from your logs. Take a few more steps, make it available to Elasticsearch and you get super-fast queries against your mountains of data.

Kibana — Visualizing Your Log Data

Kibana is your log-data dashboard. Get a better grip on your large data stores with point-and-click pie charts, bar graphs, trendlines, maps and scatter plots. You can visualize trends and patterns for data that would otherwise be extremely tedious to read and interpret. Eventually, each business line can make practical use of your data collection as you help them customize their dashboards. Save it, share it, and link your data visualizations for quick and smart communication.

How Can I Use the ELK Stack to Manage my Log Data?

Your critical business questions have answers in logs of your applications and systems, but most potential users of the data in those logs assume that the accessibility barrier is too high. But the answers are in there -- answers to questions such as:

- How many account signups this week?
- What is the effectiveness of our ad campaign?
- What is the best time to perform system maintenance?
- Why is my database performance slow?

The answers are in your data, but there are a number of challenges.

The Challenge of Disparate Log Data

If you're like most of us, your log data universe is inconsistent, inaccessible, and inconsistent. We understand problems such as these quite well:

- **No Consistency** — The variety of systems and absence of standards means that it's difficult to be a jack-of-all trades.
 - Logging is different for each app, system, or device
 - Specific knowledge is necessary for interpreting various types of logs
 - Variation in format makes it challenging to search
 - Many types of time formats
- **No centralization** — Simply put, log data is everywhere:
 - Logs in many locations on various servers

- Many locations of various logs on each server
 - SSH + GREP doesn't scale or reach
- **Accessibility of Log Data** — Much of the data is difficult to locate and manage. Although some of the log data may be highly valuable, many admins face these steep challenges:
 - Access is often difficult
 - High expertise to mine data
 - Logs can be difficult to find
 - Immense size of Log Data

The ELK stack can help you manage each of these challenges, and more. ELK is best for time-series data—anything with a timestamp—such as you'll find in most web server logs, transaction logs, and stock data listings. To be intelligible, these logs usually need substantial clean-up.

For example, you'll need to normalize in cases where you have some logs in which the timestamp is the North American dd/mm/yy, and others are the European mm/dd/yy. Logstash handles normalizations like this and many others with ease.

3. Implementation

We have deployed a scaffolded Ruby on Rails app and configured it to output its logs to the console. You will need a configured Rails environment, Docker and Docker Compose.

1. Putting the Rails app into Docker container

To get started, we need to have a working Rails app, we can use any existing app you have or create a new one:

```
$ rails new -d postgresql app
```

2. In order to support Docker environment, change the development environment in *app/config/database.yml* :

```
development: &default  
adapter: postgresql  
encoding: unicode  
database: postgres  
username: postgres  
host: db
```

3. Setup JS runtime for the app — just find the following line in the Gemfile and uncomment it:

```
gem 'therubyracer', platforms: :ruby
```

4. Add Dockerfile to the *app* folder with the following content:

```
FROM ruby:2.3  
ENV RAILS_ENV development  
WORKDIR /app  
ADD Gemfile Gemfile.lock /app/  
RUN bundle install -j5 --retry 10  
ADD . /app
```

5. Create a directory called *db* and add Dockerfile for postgres container:

```
FROM postgres  
ENV POSTGRES_USER 'postgres'  
ENV POSTGRES_DB 'app_development'
```



```
docker-compose build
Sen 8/CC Mini Project/app via v2.3.1
+ docker-compose build
Building web
Step 1/6 : FROM ruby:2.3
2.3: Pulling from library/ruby
693502eb7dfb: Already exists
081cd4bfd521: Pull complete
5d2dc01312f3: Pull complete
54a5f7da9a4f: Pull complete
168cf3f33330: Pull complete
1f2561b9c707: Pull complete
ea880d938ee1: Pull complete
d17cd70c04f3: Pull complete
Digest: sha256:064de303d9db3b01518f3ae745f92d67a3f02ceeff2c3fd10946e32a23a7b5fe
Status: Downloaded newer image for ruby:2.3
--> c61934c8d871
Step 2/6 : ENV RAILS_ENV development
--> Running in 5437b6d19c92
--> 4a59a5ee35c0
Removing intermediate container 5437b6d19c92
Step 3/6 : WORKDIR /app
--> 91be1ec69ad6
Removing intermediate container 79dd39ef5dce
Step 4/6 : ADD Gemfile Gemfile.lock /app/
--> c6f897bb5370
```

Figure 1 : Docker Compose for RAILS

```
Step 6/6 : ADD . /app
--> edf7300d1418
Removing intermediate container alafe682b5c8
Successfully built edf7300d1418
Building db
Step 1/3 : FROM postgres
latest: Pulling from library/postgres
693502eb7dfb: Already exists
31b3eb816534: Pull complete
8428acbbc47c: Pull complete
19552fac5eff: Pull complete
2f19eddaec72: Pull complete
b476232c64c2: Pull complete
e964c853634f: Pull complete
1868280d339f: Pull complete
b71b256a9a26: Pull complete
b7c0a0e9a89a: Pull complete
0c8a52608719: Pull complete
baead22f2b06: Pull complete
9d6deecab32da: Pull complete
Digest: sha256:6ac0fbeddde3bb7b0003a2ace8c126f742f8bdd90695801337d3edaaf1fcc478
Status: Downloaded newer image for postgres:latest
--> 4e18b2c30f8d
Step 2/3 : ENV POSTGRES_USER 'postgres'
--> Running in 760189904dcf
--> 540a80b16711
Removing intermediate container 760189904dcf
Step 3/3 : ENV POSTGRES_DB 'app_development'
--> Running in 1f2bb0e0a487
--> 86f8ef7299e1
Removing intermediate container 1f2bb0e0a487
Successfully built 86f8ef7299e1
```

Figure 2 : Docker Compose for RAILS

6. When the environment variables are set up, the postgres image will create a database. Create file *docker-compose.yml* in the project root directory:

```
version: '2'
services:
  web:
    build:
      context: ./app
      dockerfile: Dockerfile
    environment:
      RAILS_ENV: development
    ports:
      - '3000:3000'
    command: rails s -b 0.0.0.0

  db:
    build:
      context: ./db
      dockerfile: Dockerfile
```

7. Now, build and run the containers using the following commands:

```
$ docker-compose build
```

```
$ docker-compose up
```

As a result we are able to see our app running on <http://localhost:3000>.

```

College work/Sem 8/CC Mini Project
=> docker-compose build
Building web
Step 1/6 : FROM ruby:2.3
----> e61934e8d871
Step 2/6 : ENV RAILS_ENV development
----> Using cache
----> 4a59a5ee35c0
Step 3/6 : WORKDIR /app
----> Using cache
----> 91belec69ad6
Step 4/6 : ADD Gemfile Gemfile.lock /app/
----> Using cache
----> c6f897bb5370
Step 5/6 : RUN bundle install -j5 --retry 10
----> Using cache
----> 919e7d20668b
Step 6/6 : ADD . /app
----> 67572583b383
Removing intermediate container 86bc919013d3
Successfully built 67572583b383
Building db
Step 1/3 : FROM postgres
----> 4e18b2c30f8d
Step 2/3 : ENV POSTGRES_USER 'postgres'
----> Using cache
----> 540a80b16711
Step 3/3 : ENV POSTGRES_DB 'app_development'
----> Using cache
----> 86f8ef7299e1
Successfully built 86f8ef7299e1
elasticsearch uses an image, skipping
Building kibana

```

Figure 3 : Docker Compose for ELK

```

Status: Download newer image for kibana:latest
----> 815950a52a6e
Step 2/5 : RUN apt-get update && apt-get install -y netcat
----> Running in 0930bd7eb1d1
Get:1 https://artifacts.elastic.co stable InRelease [310 B]
Ign https://artifacts.elastic.co stable InRelease
Get:2 https://artifacts.elastic.co stable Release.gpg [473 B]
Get:3 https://artifacts.elastic.co stable Release [4360 B]
Get:4 https://artifacts.elastic.co stable/main amd64 Packages [11.0 kB]
Get:5 http://security.debian.org jessie/updates InRelease [63.1 kB]
Ign http://deb.debian.org jessie InRelease
Get:6 http://security.debian.org jessie/updates/main amd64 Packages [448 kB]
Get:7 http://deb.debian.org jessie-updates InRelease [145 kB]
Get:8 http://deb.debian.org jessie Release.gpg [2373 B]
Get:9 http://deb.debian.org jessie Release [148 kB]
Get:10 http://deb.debian.org jessie-updates/main amd64 Packages [17.6 kB]
Get:11 http://deb.debian.org jessie/main amd64 Packages [9049 kB]
Fetched 9889 kB in 16s (612 kB/s)
Reading package lists...
Reading package lists...
Building dependency tree...
Reading state information...
The following extra packages will be installed:
  netcat-traditional
The following NEW packages will be installed:
  netcat netcat-traditional
0 upgraded, 2 newly installed, 0 to remove and 1 not upgraded.
Need to get 75.3 kB of archives.
After this operation, 194 kB of additional disk space will be used.

```

Figure 4 : Docker Compose for ELK

```

Step 3/6 : COPY entrypoint.sh /tmp/entrypoint.sh
----> 4c0a8d3285d3
Removing intermediate container 6b2c28bd047d
Step 4/6 : RUN chmod +x /tmp/entrypoint.sh
----> Running in 36ee90b152c1
----> f36fcd57e72a
Removing intermediate container 36ee90b152c1
Step 5/6 : ADD kibana.yml /opt/kibana/config/
----> 24967918c81e
Removing intermediate container 9d1a93ed2476
Step 6/6 : CMD /tmp/entrypoint.sh
----> Running in 50ed12aff70b
----> 7f3b12d89b03
Removing intermediate container 50ed12aff70b
Successfully built 7f3b12d89b03
Building logstash
Step 1/2 : FROM logstash:latest
latest: Pulling from library/logstash
693502eb7dfb: Already exists
081cd4bfd521: Already exists
2e8d02ed730f: Pull complete
c5d79c1ef0b6: Pull complete
7b476804250d: Pull complete
3a3a92baa1ba: Pull complete
303c34abd3fc: Pull complete
174a19b62c3e: Pull complete
1a3d0c363309: Pull complete
53bf047e5993: Pull complete
289fa5f56b96: Pull complete
3ac19fc3d223: Pull complete
152fa8260cff: Downloading [----->] 33.95 MB/94.95 MB
328ae65e7051: Download complete
b68c30b8ecf7: Verifying Checksum

```

Figure 5 : Docker Compose for ELK

8. Setting up ELK stack

Create a folder called *logstash* in the root directory and put the Dockerfile with the following content into it:

```

FROM logstash:latest
ADD logstash.conf /etc/logstash/conf.d/

```

```

Step 6/6 : CMD /tmp/entrypoint.sh
--> Using cache
--> 7f3b12d89b03
Successfully built 7f3b12d89b03
Building logstash
Step 1/2 : FROM logstash:latest
latest: Pulling from library/logstash
693502eb7dfb: Already exists
081cd4bfd521: Already exists
2e8d00ed730f: Pull complete
c5d79e1ef0b6: Pull complete
7b476804250d: Pull complete
3a3a92baa1ba: Pull complete
303c34abd3fc: Pull complete
174a19b62c3e: Pull complete
1a3d0c363309: Pull complete
53bf047e5993: Pull complete
289fa5f56b96: Pull complete
3ac19fc3d223: Pull complete
152fa8260cff: Pull complete
328ae65e7051: Pull complete
b68c30b8ecf7: Pull complete
Digest: sha256:41963687ba3dc31b866ed03f96fa4753496ddf16271b3e2d1d1c709d43c3500c
Status: Downloaded newer image for logstash:latest
--> 413f93f5f1f2
Step 2/2 : ADD logstash.conf /etc/logstash/conf.d/
--> 659b68a6019f
Removing intermediate container 19cc8cbd4db2
Successfully built 659b68a6019f

```

Figure 6 : Building logstash

9. We also should provide a config file called *logstash.conf* (it should be placed into the same directory as our Dockerfile). We are going to consume logs from the GELF Docker log driver, so we are adding it as our input, forwarding our logs to the *elasticsearch* output, and setting up a host for it:

```

input {
  gelf {}
}

output {
  elasticsearch {
    hosts => "elasticsearch:9200"
  }
}

```

10. Create a folder called *kibana* in the root directory and put the Dockerfile with the following content into it:

```

FROM kibana:latest
RUN apt-get update && apt-get install -y netcat
COPY entrypoint.sh /tmp/entrypoint.sh
RUN chmod +x /tmp/entrypoint.sh
ADD kibana.yml /opt/kibana/config/
CMD ["/tmp/entrypoint.sh"]

```

11. Kibana should always start after the Elasticsearch so we have to add a custom entry point for our Dockerfile:

```
#!/usr/bin/env bash
while true; do
  nc -q 1 elasticsearch 9200 2>/dev/null && break
done
exec kibana
```

12. Our *kibana.yml* config file will look like this:

```
port: 5601
host: "0.0.0.0"
elasticsearch_url: "http://elasticsearch:9200"
elasticsearch_preserve_host: true
kibana_index: ".kibana"
default_app_id: "discover"
request_timeout: 300000
shard_timeout: 0
verify_ssl: true
bundled_plugin_ids:
- plugins/dashboard/index
- plugins/discover/index
- plugins/doc/index
- plugins/kibana/index
- plugins/markdown_vis/index
- plugins/metric_vis/index
- plugins/settings/index
- plugins/table_vis/index
- plugins/vis_types/index
- plugins/visualize/index
```

13. The final step is to add services with our containers to the *docker-compose.yml*:
logstash:

```
build: logstash/
command: logstash -f/etc/logstash/conf.d/logstash.conf
ports:
- "12201:12201/udp"
elasticsearch:
image: elasticsearch:latest
command: elasticsearch -Des.network.host=0.0.0.0
ports:
- "9200:9200"
- "9300:9300"
kibana:
build: kibana/
ports:
- "5601:5601"
```



```

Creating network "app_default" with the default driver
Creating app_web_1
Creating app_db_1
Attaching to app_db_1, app_web_1
db_1 | The files belonging to this database system will be owned by user "postgres".
db_1 | This user must also own the server process.
db_1 |
db_1 | The database cluster will be initialized with locale "en_US.utf8".
db_1 | The default database encoding has accordingly been set to "UTF8".
db_1 | The default text search configuration will be set to "english".
db_1 |
db_1 | Data page checksums are disabled.
db_1 |
db_1 | fixing permissions on existing directory /var/lib/postgresql/data ... ok
db_1 | creating subdirectories ... ok
db_1 | selecting default max_connections ... 100
db_1 | selecting default shared_buffers ... 128MB
db_1 | selecting dynamic shared memory implementation ... posix
db_1 | creating configuration files ... ok
db_1 | running bootstrap script ... ok
db_1 | performing post-bootstrap initialization ... ok
db_1 | syncing data to disk ...
db_1 | WARNING: enabling "trust" authentication for local connections
db_1 | You can change this by editing pg_hba.conf or using the option -A, or
db_1 | --auth-local and --auth-host, the next time you run initdb.
db_1 | ok
db_1 |
db_1 | Success. You can now start the database server using:
db_1 |
db_1 |     pg_ctl -D /var/lib/postgresql/data -l logfile start
db_1 |
db_1 | *****
db_1 | WARNING: No password has been set for the database.

```

Figure 7 : Docker Compose Up

```

db_1 | /usr/local/bin/docker-entrypoint.sh: ignoring /docker-entrypoint-initdb.d/*
db_1 |
db_1 | LOG:  received fast shutdown request
db_1 | LOG:  aborting any active transactions
db_1 | waiting for server to shut down...LOG:  autovacuum launcher shutting down
db_1 | LOG:  shutting down
db_1 | LOG:  database system is shut down
db_1 | done
db_1 | server stopped
db_1 |
db_1 | PostgreSQL init process complete; ready for start up.
db_1 |
db_1 | LOG:  database system was shut down at 2017-03-15 18:47:36 UTC
db_1 | LOG:  MultiXact member wraparound protections are now enabled
db_1 | LOG:  database system is ready to accept connections
db_1 | LOG:  autovacuum launcher started
web_1 | Started GET "/" for 172.19.0.1 at 2017-03-15 18:47:51 +0000
web_1 | Cannot render console from 172.19.0.1! Allowed networks: 127.0.0.1, ::1, 127.0.0.0/127.255.255.255
web_1 | Processing by Rails::WelcomeController#index as HTML
web_1 | Parameters: {"internal"=>true}
web_1 | Rendering /usr/local/bundle/gems/railties-5.0.2/lib/rails/templates/rails/welcome/index.html.erb
web_1 | Rendered /usr/local/bundle/gems/railties-5.0.2/lib/rails/templates/rails/welcome/index.html.erb (30.1ms)
web_1 | Completed 200 OK in 77ms (Views: 49.2ms | ActiveRecord: 0.0ms)
web_1 |
web_1 |
web_1 | ^CGracefully stopping... (press Ctrl+C again to force)
Stopping app_db_1 ... done
Stopping app_web_1 ... done

```

Figure 8 : Docker Compose Up

```

ubuntu@ip-172-26-70-43:~/Docker-ELK-Rails$ docker-compose up
Recreating dockerelkrails_elasticsearch_1
Recreating dockerelkrails_db_1
Recreating dockerelkrails_kibana_1
Recreating dockerelkrails_logstash_1
Recreating dockerelkrails_app_1
Recreating dockerelkrails_web_1
Attaching to db_postgres, elasticsearch, kibana, logstash, app_rails, nginx-proxy
db_postgres | LOG: database system was interrupted; last known up at 2017-03-19 09:12:49 UTC
db_postgres | LOG: database system was not properly shut down; automatic recovery in progress
db_postgres | LOG: invalid record length at 0/24F5AB: wanted 24, got 0
db_postgres | LOG: redo is not required
db_postgres | LOG: MultiXact member wraparound protections are now enabled
db_postgres | LOG: database system is ready to accept connections
db_postgres | LOG: autovacuum launcher started
app_rails | WARNING: no logs are available with the 'gelf' log driver
elasticsearch | [2017-03-19T13:14:22.889][INFO][e.s.n.Node] [] Initializing ...
elasticsearch | [2017-03-19T13:14:23.352][INFO][e.e.e.NodeEnvironment] [] [lifecycle] using [1] data paths, mounts [{} (overlay)], net usable_space [22.8gb], net total_space [29.3gb], spins
7 [pess10ty], types [overlay]
elasticsearch | [2017-03-19T13:14:23.364][INFO][e.e.e.NodeEnvironment] [] [lifecycle] heap size [267.6mb], compressed ordinary object pointers [true]
elasticsearch | [2017-03-19T13:14:23.403][INFO][e.e.s.Node] [] node name [lifecycle] derived from node ID [21cd0ce09a2941e28006a]; set [node.name] to override
elasticsearch | [2017-03-19T13:14:23.408][INFO][e.e.s.Node] [] version[5.2.1], pid[1], build[db6681/2017-03-09T22:05:32.386Z], 64[linux/x86_64-glibc/amd64], JVM[Oracle Corp
ration/OpenJDK 64-Bit Server VM/1.8.0_92-internal/25.92-b14]
elasticsearch | [2017-03-19T13:14:23.466][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [aggs-matrix-stats]
elasticsearch | [2017-03-19T13:14:23.478][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [ingest-common]
elasticsearch | [2017-03-19T13:14:23.478][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [lang-expression]
elasticsearch | [2017-03-19T13:14:23.478][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [lang-groovy]
elasticsearch | [2017-03-19T13:14:23.479][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [lang-mustache]
elasticsearch | [2017-03-19T13:14:23.479][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [lang-painless]
elasticsearch | [2017-03-19T13:14:23.483][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [percolator]
elasticsearch | [2017-03-19T13:14:23.483][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [reindex]
elasticsearch | [2017-03-19T13:14:23.483][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [transport-netty3]
elasticsearch | [2017-03-19T13:14:23.483][INFO][e.e.p.PluginsService] [] [lifecycle] loaded module [transport-netty4]
elasticsearch | [2017-03-19T13:14:23.485][INFO][e.e.p.PluginsService] [] [lifecycle] loaded plugin [analysis]

```

Figure 9 : Docker Compose Up

```

ubuntu@ip-172-26-70-43:~/Docker-ELK-Rails$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS               NAMES
4292277a0e79       dockerelkrails_web  "nginx -g 'daemon ..." 13 minutes ago      Up 13 minutes      0.0.0.0:80->80/tcp, 443/tcp    nginx-proxy
3406de2276de       dockerelkrails_app  "bundle exec rails..." 13 minutes ago      Up 13 minutes      0.0.0.0:3000->3000/tcp        app_rails
7f33004db19d       dockerelkrails_logstash  "logstash -f /etc/..." 13 minutes ago      Up 13 minutes      0.0.0.0:12201->12201/udp      logstash
5810e0d0fb8       dockerelkrails_db      "docker-entrypoint..." 13 minutes ago      Up 13 minutes      5432/tcp                  db_postgres

ubuntu@ip-172-26-70-43:~/Docker-ELK-Rails$ docker-compose stop
Stopping nginx-proxy ... done
Stopping app_rails ... done
Stopping logstash ...
Stopping db_postgres ... done

```

Figure 10 : Docker ps

14. After building and running the services we'll be able to access Kibana on <http://localhost:5601>. If we want the logs to be persistent — we just need to put our Elasticsearch data folder (`/usr/share/elasticsearch/data`) to a volume.

15. Sending Rails logs to Logstash

In order to forward our container logs to logstash we should just turn on GELF driver for our `web` container and set up a host URL:

`web:`

`...`

`logging:`

`driver: gelf`

`options:`

`gelf-address: 'udp://localhost:12201'`

Note: why do we use `localhost` instead of `logstash` here? The code which configures logging driver will be executed on the host machine which is not part of default network created by docker.

16. After restarting our services, open the Rails app, open Kibana, navigate to the settings tab and create a index:

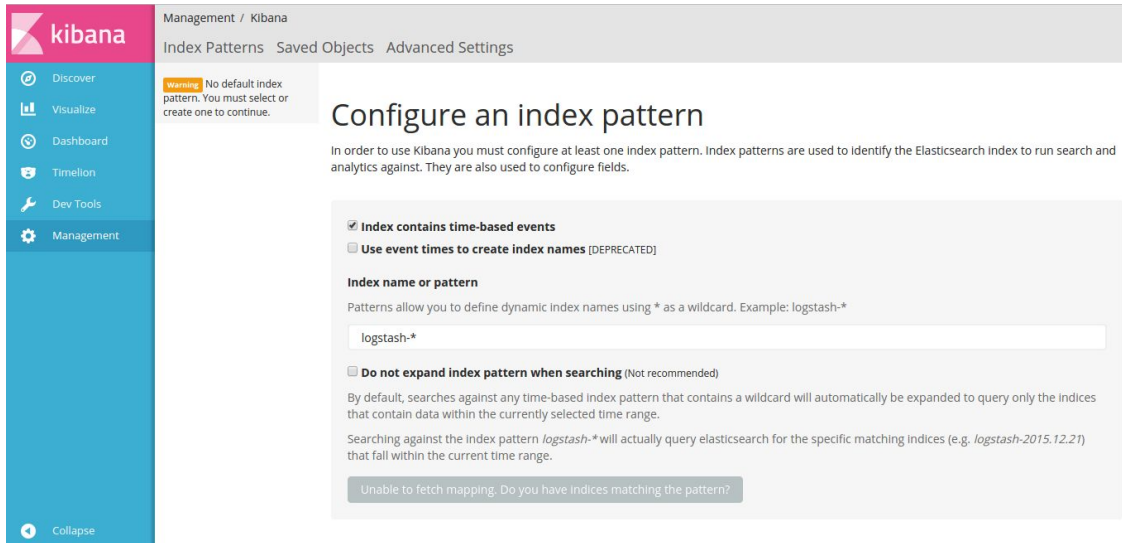


Figure 11 : Creating an index in Kibana

When index is created, we'll be able to see our first log entries in the Kibana UI. However, our current logging is not super helpful since we just sending plaintext from the Rails log:

@timestamp	August 30th 2016, 18:39:07.036
@version	1
_id	AVbcGq8sa3c9j8Eqk2k3
_index	logstash-2016.08.30
_score	
_type	logs
command	rails s -b 0.0.0.0
container_id	e63328242d48595e94358a31e1883764d383cc247792595aee5ce0cb8ddb4e3a
container_name	elkdemo_web_1
created	August 30th 2016, 18:37:07.575
host	moby
image_id	sha256:4f2d647b442f2f7a0c8676b87c43c40eebbf128f46c80b7e7e183201e3c4e66d
image_name	elkdemo_web
level	6
message	{"method":"GET","path":"/","format":"html","controller":"Rails::WelcomeController","action":"index","status":200,"duration":14.59,"view":7.07,"db":0.0,"request_ip":null,"@timestamp":"2016-08-30T15:39:07.035Z","@version":"1","message":["[200] GET / (Rails::WelcomeController#index)"]}
source_host	172.24.0.1
tag	
version	1.1

Figure 12 : Log entry

t controller	🔍 🔍 📄 Rails::WelcomeController
🕒 created	🔍 🔍 📄 August 27th 2016, 16:05:59.787
# db	🔍 🔍 📄 0
# duration	🔍 🔍 📄 14.68
t format	🔍 🔍 📄 html
t host	🔍 🔍 📄 moby
t image_id	🔍 🔍 📄 sha256:cd7a85140b8d8a9d9e05d715580a69d154509c535d9b53fd79a3ba6d8c0e6497
t image_name	🔍 🔍 📄 e1kdemo_web
# level	🔍 🔍 📄 6
t method	🔍 🔍 📄 GET
t path	🔍 🔍 📄 /
? request_ip	🔍 🔍 📄 ⚠️ -
t source_host	🔍 🔍 📄 172.24.0.1
# status	🔍 🔍 📄 200

Figure 13 : Log entry

17. Instead we should present our log entries in JSON format and put as much useful data as we can. We shall use the [Lograge](#) gem, we should add two gems to our Gemfile: *lograge* and *logstash-event*. After that we should create a new initializer called *lograge.rb*, which will set up an appropriate format for our logs and send them to stdout where docker will grab them and send to the GELF driver:

```
Rails.application.configure do
  config.lograge.formatter = Lograge::Formatters::Logstash.new
  config.lograge.logger = ActiveSupport::Logger.new(STDOUT)
End
```

18. In order to turn on lograge for the environment (now we want it for the *development.rb*) — just add the following line to the corresponding environment file:
 config.lograge.enabled = true

19. Adding custom fields is very simple, if we need something from the controller context — just overload the *append_info_to_payload* method in the ApplicationController and put everything you need to the payload, and after that we'll be able to add it to log. For instance, if we want to send an IP where the request came from, your *append_info_to_payload* will be:

```
class ApplicationController < ActionController::Base
  ...
  def append_info_to_payload(payload)
    super
    payload[:request_ip] = request.ip
  end
End
```

And *lograge.rb* should contain:

```

config.lograge.custom_options = lambda do |event|
  {
    request_ip: event.payload[:request_ip]
  }
end

```

20. The final step is to modify our *logstash.conf* file to parse the incoming JSON from the *message* field, just add the following code after the *input* section:

```

filter {
  json {
    source => "message"
    remove_field => "message"
  }
}

```

After rebuilding and starting your services we'll see that all the data coming from the Rails container is structured and indexed by Elasticsearch:

21. Search for all the GET requests : We should open the *Discover* tab and type a query "method:GET" to the search field. We'll see all the log entries with the GET HTTP method with the highlights:



The screenshot shows the Elasticsearch Discover tab with two log entries. Each entry is a JSON object containing metadata and log data. The first entry is from August 30th, 2016, at 23:40:30.942, and the second is from August 30th, 2016, at 23:41:04.233. Both entries are for GET requests to the Rails application, with the controller being Rails::WelcomeController. The log data is highlighted in yellow.

```

> August 30th 2016, 23:40:30.942 method: GET version: 1.1 host: moby level: 6 @version: 1 @timestamp: August 30th 2016, 23:41:17.321 source_host: 1
72.24.0.1 command: rails s -b 0.0.0.0 container_id: b4c093e3528e1560b0c50691c54f7c2954176864d695c8424af7f4d6791b0f79
container_name: elkdemo_web_1 created: August 30th 2016, 23:40:30.942 image_id: sha256:a4b44a15fd1387e261548a34179f2f3d
707339c4c9f7d25557227d52147b9906 image_name: elkdemo_web tag: path: / format: html controller: Rails::WelcomeControl
ter action: index status: 200 duration: 5.8 view: 3.89 db: 0 request ip: - id: AVbdl0Vv143Ez1E133Vn type: logs

> August 30th 2016, 23:40:30.942 method: GET version: 1.1 host: moby level: 6 @version: 1 @timestamp: August 30th 2016, 23:41:04.233 source_host: 1
72.24.0.1 command: rails s -b 0.0.0.0 container_id: b4c093e3528e1560b0c50691c54f7c2954176864d695c8424af7f4d6791b0f79
container_name: elkdemo_web_1 created: August 30th 2016, 23:40:30.942 image_id: sha256:a4b44a15fd1387e261548a34179f2f3d
707339c4c9f7d25557227d52147b9906 image_name: elkdemo_web tag: path: / format: html controller: Rails::WelcomeControl
ter action: index status: 200 duration: 5.66 view: 2.92 db: 0 request ip: - id: AVbdlx1v143Ez1E133Vn type: logs

```

Figure 14 : Log entry in JSON format

4. Results and Graphs

Visualization of Log Entries using Bar Graphs

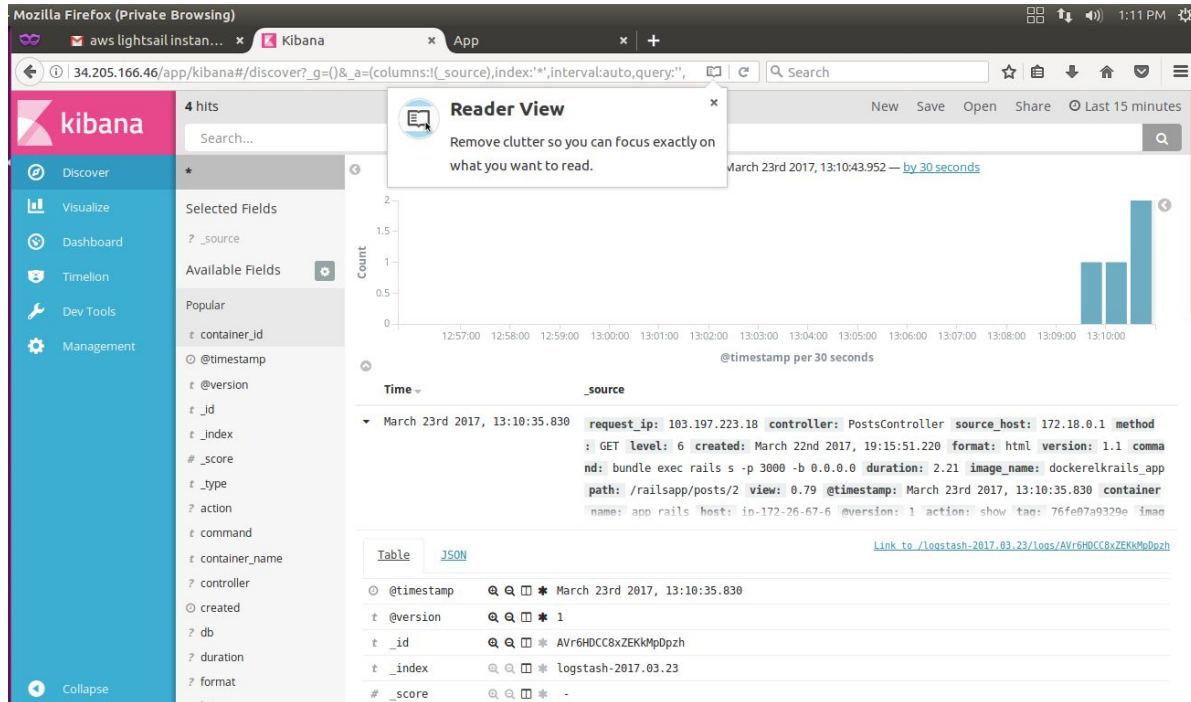


Figure 15 : Bar Graph showing number of requests per day.

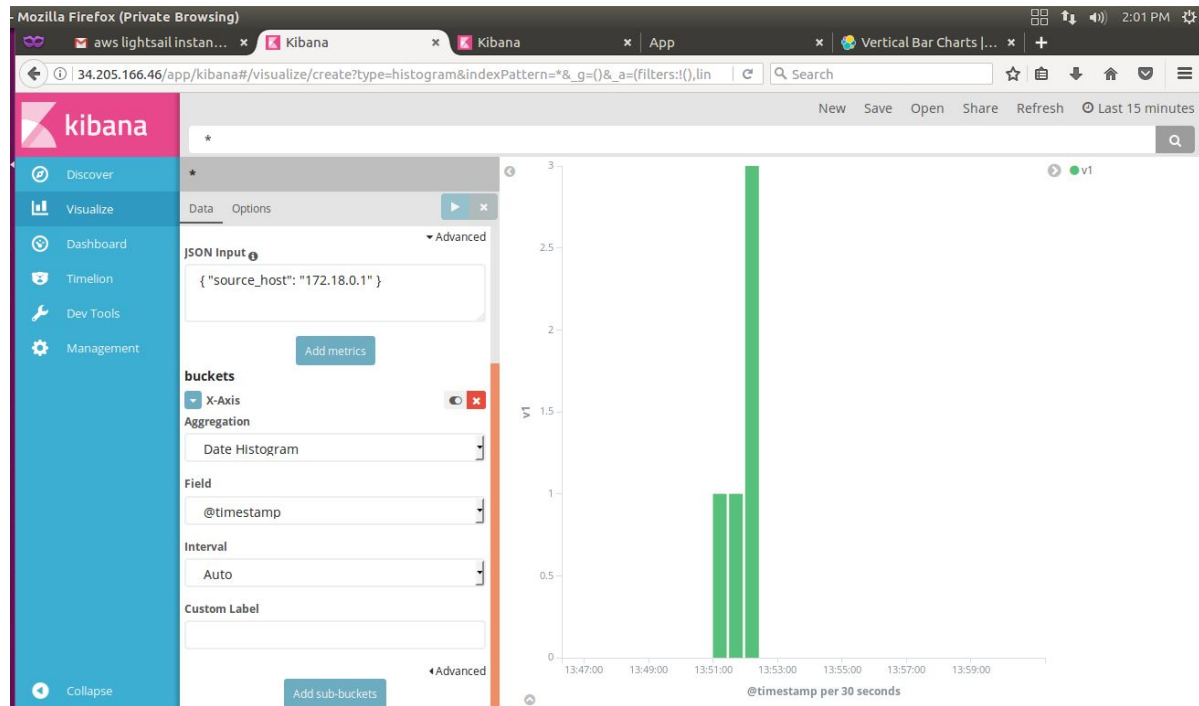


Figure 17 : Bar Graphs depicting number of requests from ip=172.18.0.1 per 30 seconds.

5. Conclusion

User can access our Rails app on the <IP>/railsapp/posts and Kibana on <IP>/. The Ruby on Rails app sends logs to Logstash via 12201/udp port (GELF input endpoint) and Kibana and Logstash talk to Elasticsearch via port 9200 (JSON REST API port). Nginx serves as a reverse-proxy server and web server and routes requests to Kibana or the Rails app based on the endpoint that you visit.

The application was successfully deployed on AWS Lightsail server with 2GB memory, 1 CPU and 40GB SSD.

[

6. References

- Docker Monitoring with the ELK Stack: A Step-by-Step Guide
<https://logz.io/learn/docker-monitoring-elk-stack/>
- How to setup centralized logging on your Django apps using Logstash and Amazon EC2
<https://csparpa.github.io/blog/2016/10/django-centralized-logging-on-ec2-with-logstash.html>
- How to set up ELK for Rails log management using Docker and Docker Compose
<https://medium.com/@AnjLab/how-to-set-up-elk-for-rails-log-management-using-docker-and-docker-compose-a6edc290669f>